

Validación métricas con HiperParametros - Laboratorio 3

Luis Felipe Sanchez Sanchez - 90613
Harold Shneider Martinez Tapiero - 77999

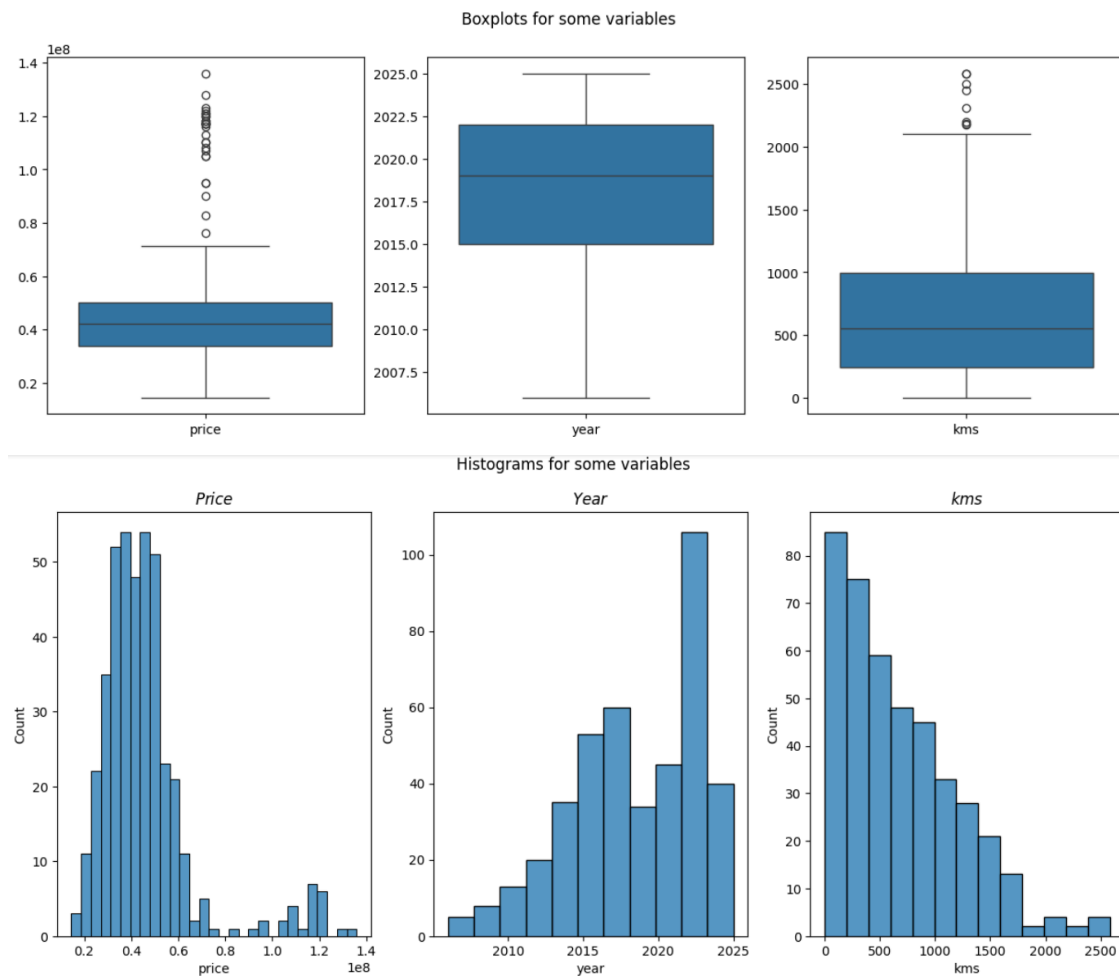
Universidad ECCI
Facultad de Ingeniería

Elias Buitrago Bolivar

Bogota
2024

Resultados de los diferentes modelos con los datos subidos sin seleccionar ni filtrar ninguna info:

Graficas:



Multivariate lineal regression

```
✓ [87] 1 # Define model and prediction
      2 ols = LinearRegression()
      3 model1 = ols.fit(X_train, y_train)
      4 y_pred1 = model1.predict(X_test)

✓ [88] 1 # accuracy check
      2 rmse = MSE(y_test, y_pred1, squared=False)
      3 mae = MAE(y_test, y_pred1)
      4 r2 = r2_score(y_test, y_pred1)
      5 print("RMSE: %.2f" % rmse)
      6 print("MAE: %.2f" % mae)
      7 print("R2: %.2f" % r2)

RMSE: 12109549.03
MAE: 8817707.26
R2: 0.57
```

RMSE: 12109549.03

MAE: 8817707.26

R2: 0.57

Light GBM

```
[91] 1 # Hyperparameters
2 params = {
3     'task': 'train',
4     'boosting': 'gbdt',
5     'objective': 'regression',
6     'num_leaves': 10,
7     'learning_rate': 0.05,
8     'metric': {'l2', 'l1'},
9     'header': 'true',
10    'verbose': 0
11 }
12
13 # loading data
14 lgb_train = lgb.Dataset(X_train, y_train)
15 lgb_eval = lgb.Dataset(X_test, y_test, reference=lgb_train)
16
17 # fitting the model
18 model2 = lgb.train(params,
19                    train_set=lgb_train,
20                    valid_sets=lgb_eval)
21 # Pred
22 y_pred2 = model2.predict(X_test)

[92] 1 # accuracy check
2 rmse = MSE(y_test, y_pred2, squared=False)
3 mae = MAE(y_test, y_pred2)
4 r2 = r2_score(y_test, y_pred2)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)

RMSE: 11126993.07
MAE: 6700310.15
R2: 0.63
```

RMSE: 11126993.07

MAE: 6700310.15

R2: 0.63

Random Forest Regressor

```
Random Forest Regressor

[96] 1 from sklearn.ensemble import RandomForestRegressor

[97] 1 model3 = RandomForestRegressor()
2 model3.fit(X_train, y_train)
3 y_pred3 = model3.predict(X_test)

[98] 1 # accuracy check
2 rmse = MSE(y_test, y_pred3, squared=False)
3 mae = MAE(y_test, y_pred3)
4 r2 = r2_score(y_test, y_pred3)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)

RMSE: 9477736.16
MAE: 5146986.60
R2: 0.73
```

RMSE: 9477736.16

MAE: 5146986.60

R2: 0.73

Xgboost regressor

```
12s [103] 1 #K-fold cross validation
2 scores = cross_val_score(model4, X_train, y_train, cv=10)
3 print("Mean cross-validation score: %.2f" % scores.mean())

Mean cross-validation score: 0.68

12s [104] 1 kfold = KFold(n_splits=10, shuffle=True)
2 kf_cv_scores = cross_val_score(model4, X_train, y_train, cv=kfold)
3 print("K-fold CV average score: %.2f" % kf_cv_scores.mean())

K-fold CV average score: 0.71

0s [105] 1 # Pred
2 y_pred4 = model4.predict(X_test)

0s [106] 1 # accuracy check
2 rmse = MSE(y_test, y_pred4, squared=False)
3 mae = MAE(y_test, y_pred4)
4 r2 = r2_score(y_test, y_pred4)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)

RMSE: 10238249.56
MAE: 5221331.12
R2: 0.69
```

RMSE: 10238249.56

MAE: 5221331.12

R2: 0.69

Prueba 1:

Light GBM

```
0s [107] 1 # Hyperparameters
2 params = {
3     'task': 'train',
4     'boosting': 'gbdt',
5     'objective': 'regression',
6     'num_leaves': 9,
7     'learning_rate': 0.21,
8     'metric': {'12', '11'},
9     'header': 'true',
10    'verbose': 0
11 }
12
13 # loading data
14 lgb_train = lgb.Dataset(X_train, y_train)
15 lgb_eval = lgb.Dataset(X_test, y_test, reference=lgb_train)
16
17 # fitting the model
18 model2 = lgb.train(params,
19                    train_set=lgb_train,
20                    valid_sets=lgb_eval)
21 # Pred
22 y_pred2 = model2.predict(X_test)

0s [108] 1 # accuracy check
2 rmse = MSE(y_test, y_pred2, squared=False)
3 mae = MAE(y_test, y_pred2)
4 r2 = r2_score(y_test, y_pred2)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)

RMSE: 12625286.83
MAE: 6797070.60
R2: 0.70
```

RMSE: 12625286.83

MAE: 6797070.60

R2: 0.70

Random Forest Regressor

```
[ ] 1 from sklearn.ensemble import RandomForestRegressor
```

```
[ ] 1 model3 = RandomForestRegressor()  
2 model3.fit(X_train, y_train)  
3 y_pred3 = model3.predict(X_test)
```

```
[ ] 1 # accuracy check  
2 rmse = MSE(y_test, y_pred3, squared=False)  
3 mae = MAE(y_test, y_pred3)  
4 r2 = r2_score(y_test, y_pred3)  
5 print("RMSE: %.2f" % rmse)  
6 print("MAE: %.2f" % mae)  
7 print("R2: %.2f" % r2)
```



RMSE: 8993766.98
MAE: 5157716.40
R2: 0.85

RMSE: 8993766.98

MAE: 5157716.40

R2: 0.85

Xgboost regressor

```
[ ] 1 # Pred  
2 y_pred4 = model4.predict(X_test)
```

```
[ ] 1 # accuracy check  
2 rmse = MSE(y_test, y_pred4, squared=False)  
3 mae = MAE(y_test, y_pred4)  
4 r2 = r2_score(y_test, y_pred4)  
5 print("RMSE: %.2f" % rmse)  
6 print("MAE: %.2f" % mae)  
7 print("R2: %.2f" % r2)
```



RMSE: 11535189.22
MAE: 6249334.76
R2: 0.75

RMSE: 11535189.22

MAE: 6249334.76

R2: 0.75

Prueba 2:

Light GBM

```
[80] 1 # Hyperparameters
2     params = {
3         'task': 'train',
4         'boosting': 'gbdt',
5         'objective': 'regression',
6         'num_leaves': 5,
7         'learning_rate': 0.52,
8         'metric': {'l2', 'l1'},
9         'header': 'true',
10        'verbose': 0
11    }
12
13 # loading data
14 lgb_train = lgb.Dataset(X_train, y_train)
15 lgb_eval = lgb.Dataset(X_test, y_test, reference=lgb_train)
16
17 # fitting the model
18 model2 = lgb.train(params,
19                    train_set=lgb_train,
20                    valid_sets=lgb_eval)
21 # Pred
22 y_pred2 = model2.predict(X_test)
```

```
1 # accuracy check
2 rmse = MSE(y_test, y_pred2, squared=False)
3 mae = MAE(y_test, y_pred2)
4 r2 = r2_score(y_test, y_pred2)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)
```

RMSE: 11861025.88
MAE: 6673291.02
R2: 0.74

RMSE: 11861025.88

MAE: 6673291.02

R2: 0.74

Random Forest Regressor

```
1 param_grid = {
2     'n_estimators': [50, 100, 200, 300], # Número de árboles en
3     'max_depth': [None, 10, 20, 30, 40, 50], # Profundidad máxi
4     'min_samples_split': [2, 5, 10], # Número mínimo de muestra
5     'min_samples_leaf': [1, 2, 4], # Número mínimo de muestras
6     'max_features': ['auto', 'sqrt', 'log2'] # Número máximo de
7 }
8 rf = RandomForestRegressor(random_state=42)
9 random_search = RandomizedSearchCV(estimator=rf, param_distribut
10                                     scoring=scoring, refit='rmse'
11
12
13 random_search.fit(X_train, y_train)
14 best_model = random_search.best_estimator_
15 y_pred = best_model.predict(X_test)
16
17
18 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
19 mae = mean_absolute_error(y_test, y_pred)
20 r2 = r2_score(y_test, y_pred)
21
22 print("RMSE: %.2f" % rmse)
23 print("MAE: %.2f" % mae)
24 print("R2: %.2f" % r2)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits
RMSE: 8900388.53
MAE: 5350869.41
R2: 0.83

RMSE: 8900388.53

MAE: 5350869.41

R2: 0.83

Xgboost regressor

```
1 # accuracy check
2 rmse = MSE(y_test, y_pred4, squared=False)
3 mae = MAE(y_test, y_pred4)
4 r2 = r2_score(y_test, y_pred4)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)
```

```
RMSE: 12371077.69
MAE: 7126827.66
R2: 0.67
```

RMSE: 12371077.69

MAE: 7126827.66

R2: 0.67

Prueba 3:

Light GBM

```
1 # Hyperparameters
2 params = {
3     'task': 'train',
4     'boosting': 'gbdt',
5     'objective': 'regression',
6     'num_leaves': 5,
7     'learning_rate': 1.0,
8     'metric': {'l2', 'l1'},
9     'header': 'true',
10    'verbose': 0
11 }
12
13 # loading data
14 lgb_train = lgb.Dataset(X_train, y_train)
15 lgb_eval = lgb.Dataset(X_test, y_test, reference=lgb_train)
16
17 # fitting the model
18 model2 = lgb.train(params,
19                    train_set=lgb_train,
20                    valid_sets=lgb_eval)
21
22 # Pred
23 y_pred2 = model2.predict(X_test)
```

```
1 # accuracy check
2 rmse = MSE(y_test, y_pred2, squared=False)
3 mae = MAE(y_test, y_pred2)
4 r2 = r2_score(y_test, y_pred2)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)
```

```
RMSE: 11597556.08
MAE: 6994310.87
R2: 0.75
```

RMSE: 11597556.08
MAE: 6994310.87
R2: 0.75

Random Forest Regressor

```
1 param_grid = {
2     'n_estimators': [ 150, 200, 350], # Número de árboles en el
3     'max_depth': [None, 10, 20], # Profundidad máxima de los ár
4     'min_samples_split': [2, 9, 10], # Número mínimo de muestr
5     'min_samples_leaf': [2, 3, 4], # Número mínimo de muestras
6     'max_features': ['auto', 'sqrt', 'log2'] # Número máximo de
7 }
8 rf = RandomForestRegressor(random_state=42)
9 random_search = RandomizedSearchCV(estimator=rf, param_distribut
10                                     scoring=scoring, refit='rmse'
11
12
13 random_search.fit(X_train, y_train)
14 best_model = random_search.best_estimator_
15 y_pred = best_model.predict(X_test)
16
17
18 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
19 mae = mean_absolute_error(y_test, y_pred)
20 r2 = r2_score(y_test, y_pred)
21
22 print("RMSE: %.2f" % rmse)
23 print("MAE: %.2f" % mae)
24 print("R2: %.2f" % r2)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits
RMSE: 8684526.04
MAE: 5170397.83
R2: 0.84

RMSE: 8684526.04
MAE: 5170397.83
R2: 0.84

Xgboost regressor


```
[50] 1 # Pred
      2 y_pred4 = model4.predict(X_test)
```

```
1 # accuracy check
2 rmse = MSE(y_test, y_pred4, squared=False)
3 mae = MAE(y_test, y_pred4)
4 r2 = r2_score(y_test, y_pred4)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)
```

```
RMSE: 11715237.49
MAE: 6850894.90
R2: 0.71
```

RMSE: 11715237.49

MAE: 6850894.90

R2: 0.71

Prueba 4:

Light GBM

```
12
13 # loading data
14 lgb_train = lgb.Dataset(X_train, y_train)
15 lgb_eval = lgb.Dataset(X_test, y_test, reference=lgb_train)
16
17 # fitting the model
18 model2 = lgb.train(params,
19                    train_set=lgb_train,
20                    valid_sets=lgb_eval)
21 # Pred
22 y_pred2 = model2.predict(X_test)
```

```
1 # accuracy check
2 rmse = MSE(y_test, y_pred2, squared=False)
3 mae = MAE(y_test, y_pred2)
4 r2 = r2_score(y_test, y_pred2)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)
```

```
RMSE: 11228274.90
MAE: 6582336.57
R2: 0.76
```

RMSE: 11228274.90

MAE: 6582336.57

R2: 0.76

Random Forest Regressor

```
1 param_grid = {
2     'n_estimators': [ 200, 400], # Número de árboles en el bo
3     'max_depth': [None, 10, 20, 30], # Profundidad máxima de
4     'min_samples_split': [2, 9, 10], # Número mínimo de mues
5     'min_samples_leaf': [2, 3, 4, 5], # Número mínimo de mues
6     'max_features': ['auto', 'sqrt', 'log2'] # Número máximo
7 }
8 rf = RandomForestRegressor(random_state=42)
9 random_search = RandomizedSearchCV(estimator=rf, param_distrib
10                                     scoring=scoring, refit='rms
11
12
13 random_search.fit(X_train, y_train)
14 best_model = random_search.best_estimator_
15 y_pred = best_model.predict(X_test)
16
17
18 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
19 mae = mean_absolute_error(y_test, y_pred)
20 r2 = r2_score(y_test, y_pred)
21
22 print("RMSE: %.2f" % rmse)
23 print("MAE: %.2f" % mae)
24 print("R2: %.2f" % r2)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits
RMSE: 8955562.17
MAE: 5367564.00
R2: 0.83

RMSE: 8955562.17

MAE: 5367564.00

R2: 0.83

Xgboost regressor

```
[63] 1 # Pred
      2 y_pred4 = model4.predict(X_test)
```

```
1 # accuracy check
2 rmse = MSE(y_test, y_pred4, squared=False)
3 mae = MAE(y_test, y_pred4)
4 r2 = r2_score(y_test, y_pred4)
5 print("RMSE: %.2f" % rmse)
6 print("MAE: %.2f" % mae)
7 print("R2: %.2f" % r2)
```

RMSE: 10775624.04
MAE: 6505791.87
R2: 0.75

RMSE: 10775624.04

MAE: 6505791.87

R2: 0.75